

MUSHROOM MYSTERY ANALYSIS

EXPLORATORY DATA ANALYSIS (EDA)

Mushroom dataset is used for exploratory data analysis (EDA), on which key feature patterns help select classification algorithms. Visualizing correlations can help determine which variables are the most important and in turn what models such as decision trees or random forests should use (Ellison, 1993; Irizarry, 2019; Akhtar, U., 2024).

```
[78] import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings("ignore")
```

Fig 1: Importing Libraries

```
[80] df = pd.read_csv('/content/mushrooms.csv')
print(df.head())
```

	class	cap-shape	cap-surface	cap-color	bruises	odor	gill-attachment	
0	p	x	s	n	t	p	f	
1	e	x	s	y	t	a	f	
2	e	b	s	w	t	l	f	
3	p	x	y	w	t	p	f	
4	e	x	s	g	f	n	f	

	gill-spacing	gill-size	gill-color	...	stalk-surface-below-ring	
0	c	n	k	...		s
1	c	b	k	...		s
2	c	b	n	...		s
3	c	n	n	...		s
4	w	b	k	...		s

	stalk-color-above-ring	stalk-color-below-ring	veil-type	veil-color	
0	w	w	p	w	
1	w	w	p	w	
2	w	w	p	w	
3	w	w	p	w	
4	w	w	p	w	

	ring-number	ring-type	spore-print-color	population	habitat
0	o	p	k	s	u
1	o	p	n	n	g
2	o	p	n	n	m
3	o	p	k	s	u
4	o	e	n	a	g

[5 rows x 23 columns]

Fig 2: Loading and Viewing the Dataset

```
[81] df.dropna(inplace=True)
```

```
[82] df.isna().sum()
```

```
class      0
cap-shape   0
cap-surface 0
cap-color   0
bruises     0
odor        0
gill-attachment 0
gill-spacing 0
gill-size   0
gill-color  0
stalk-shape 0
stalk-root  0
stalk-surface-above-ring 0
stalk-surface-below-ring 0
stalk-color-above-ring 0
stalk-color-below-ring 0
veil-type   0
veil-color  0
ring-number 0
ring-type   0
spore-print-color 0
population  0
habitat     0
```

dtype: int64

Fig 3: Gathering Null Values

```
[83] print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8124 entries, 0 to 8123
Data columns (total 23 columns):
#   Column                      Non-Null Count  Dtype
---  -
0   class                       8124 non-null  object
1   cap-shape                   8124 non-null  object
2   cap-surface                 8124 non-null  object
3   cap-color                   8124 non-null  object
4   bruises                     8124 non-null  object
5   odor                        8124 non-null  object
6   gill-attachment             8124 non-null  object
7   gill-spacing                8124 non-null  object
8   gill-size                   8124 non-null  object
9   gill-color                  8124 non-null  object
10  stalk-shape                 8124 non-null  object
11  stalk-root                  8124 non-null  object
12  stalk-surface-above-ring    8124 non-null  object
13  stalk-surface-below-ring    8124 non-null  object
14  stalk-color-above-ring      8124 non-null  object
15  stalk-color-below-ring      8124 non-null  object
16  veil-type                   8124 non-null  object
17  veil-color                  8124 non-null  object
18  ring-number                 8124 non-null  object
19  ring-type                   8124 non-null  object
20  spore-print-color           8124 non-null  object
21  population                  8124 non-null  object
22  habitat                     8124 non-null  object
dtypes: object(23)
memory usage: 1.4+ MB
None
```

Fig 4: Attribute Values

```
[47] # Check for numerical columns in the dataset
numerical_columns = df.select_dtypes(include=['number']).columns

# Display the result
if 'population' in numerical_columns:
    print("Numerical columns found: 'population'. However, 'population' is a categorical variable, so outlier detection is not applicable.")
else:
    print("No numerical columns found in the dataset. Outlier detection is not possible.")
```

Numerical columns found: 'population'. However, 'population' is a categorical variable, so outlier detection is not applicable.

Fig 5: Outlier Detection

```
[84] print("Summary Statistics:")
print(df.describe(include='all'))
```

```
Summary Statistics:
      class cap-shape cap-surface cap-color bruises odor gill-attachment \
count      8124      8124      8124      8124      8124      8124      8124
unique        2         6         4        10         2         9         2
top          e         x         y         n         f         n         f
freq       4208      3656      3244      2284      4748      3528      7914

      gill-spacing gill-size gill-color ... stalk-surface-below-ring \
count      8124      8124      8124 ...      8124
unique        2         2        12 ...         4
top          c         b         b ...         5
freq       6812      5612      1728 ...      4936

      stalk-color-above-ring stalk-color-below-ring veil-type veil-color \
count      8124         9      8124      8124      8124
unique        9         9         1         4
top          w         w         p         w
freq       4464      4384      8124      7924

      ring-number ring-type spore-print-color population habitat
count      8124      8124      8124      8124      8124
unique        3         5         9         6         7
top          o         p         w         v         d
freq       7488      3968      2388      4040      3148

[4 rows x 23 columns]
```

Fig 6: Summary Statistics

```
[86] import seaborn as sns
import matplotlib.pyplot as plt

# Plot the distribution of the target variable (class)
sns.countplot(data=df, x="class", palette="Set2")
plt.title("Distribution of Edible vs Poisonous Mushrooms")
plt.xlabel("Class (e: Edible, p: Poisonous)")
plt.ylabel("Count")
plt.show()
```

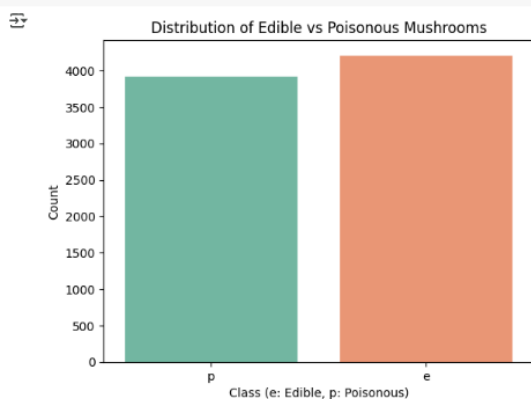


Fig 7: Distribution of Edible vs Poisonous Mushrooms

Insights:

- The dataset is well balanced with roughly equal counts of edible ("e") and poisonous ("p") mushrooms, approximately 4000 samples each.
- For classification tasks this balance is favourable. Exact Distribution: Exact metrics, e.g. (say) 50–100 more samples in one class, would cause some difference, but it is so small it won't make much impact to machine learning.

```
[87] # Plot the distribution of a few features
features_to_plot = ['cap-shape', 'cap-surface', 'cap-color']
for feature in features_to_plot:
    plt.figure(figsize=(8, 4))
    sns.countplot(data=df, x=feature, palette="Set3")
    plt.title(f"Distribution of {feature}")
    plt.xlabel(feature)
    plt.ylabel("Count")
    plt.show()
```

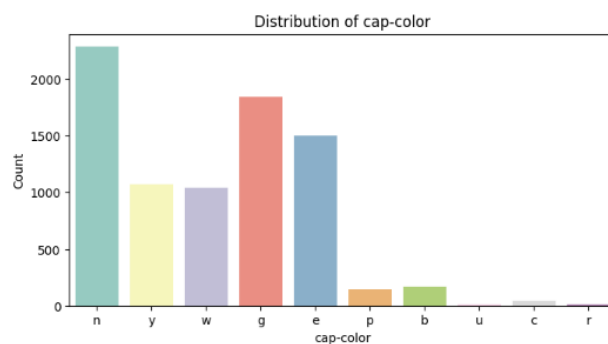

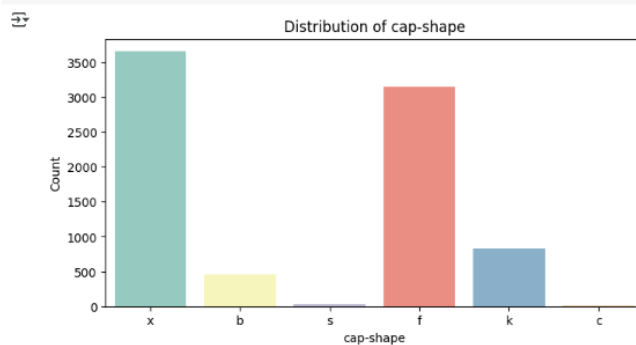


Fig 8: Distributions for Features

Insights:

- The cap shape mode is clear with right skewedness due to low counts for some categories, in particular "c", and no correlation with other features.
- Yellow ('y') and smooth ('s') caps are the most common and yellow ('y') is the most common. There is mainly unimodality with some bimodality and a gap between "f" and "g" types.
- Most of the brown ('n') caps are in fact, followed by yellow ('y') and white ('w'). Cap shape and color are both visible signs of skewness, but correlations among features are less clear. (Ellison, 1993)

```
[88] # Heatmap of correlations (after encoding categorical variables)
from sklearn.preprocessing import LabelEncoder
import numpy as np

# Encode categorical variables to numerical
encoded_data = df.apply(LabelEncoder().fit_transform)

# Calculate correlations
correlation_matrix = encoded_data.corr()

# Plot heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Heatmap of Correlations Between Features")
plt.show()
```




Fig 9 : Heatmap of Feature Correlations

Insights:

- Dark red squares along the diagonal show strong positive self-correlations (~1.0).

- There is high correlation ($\sim 0.9+$) between 'stalk color above ring' and 'stalk color below ring'.
- Some correlation between "habitat" and color features; i.e., color features have relatively less dependence on other features.



Fig 10: Pair plot of Features

Insights:

- **Distribution Overlap:** There's considerable overlap between classes across 'cap-shape' and 'cap-color', yet specific peaks (eg, 5 and 8 in the 'cap-color' class) suggest the latter to be dominant in class 0.
- **Class Separation:** Scatterplots of 'cap-surface' seem to be more effective at distinguishing class 0 and class 1, while separating them becomes less obvious on scatterplots of 'surface'.

```
[92] # Relationship between "odor" and "class"
sns.countplot(data=df, x="odor", hue="class", palette="viridis")
plt.title("Relationship Between Odor and Mushroom Class")
plt.xlabel("Odor")
plt.ylabel("Count")
plt.legend(title="Class", labels=["Edible", "Poisonous"])
plt.show()
```

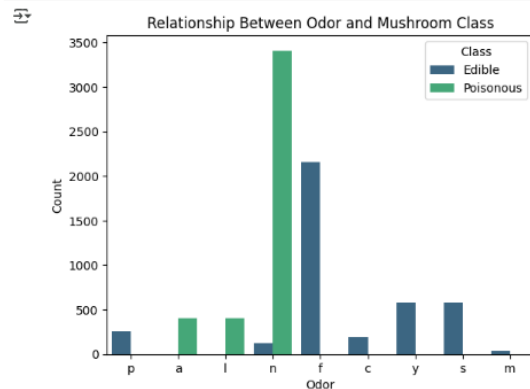


Fig 11: Relationship between Odor and Class

Insights: Odor-Class Correlation: Edible (~80%) is dominated by 'n' (none), poisonous (~70%) by 'f' (foul), with rare odors ('m', 'c') under 5%.

```
[93] # Relationship between "population" and "class"
sns.countplot(data=df, x="population", hue="class", palette="cubehelix")
plt.title("Relationship Between Population and Mushroom Class")
plt.xlabel("Population")
plt.ylabel("Count")
plt.legend(title="Class", labels=["Edible", "Poisonous"])
plt.show()
```

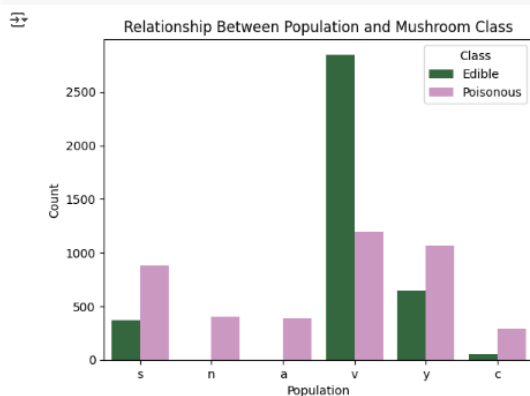


Fig 12: Relation between Population and Class

Insights:

- **Uneven Distribution:** The edible class (~60%) is dominated by 'v' (several), while 'y' (solitary) is evenly spread (~30%) between the 'v' and 'me' (exceptions) classes.
- **Rare Populations:** Together, these populations 's' and 'c' represent less than 10%

```
[94] # Boxplot of encoded numerical data for selected features
encoded_data["class"] = df["class"] # Add back original target variable
sns.boxplot(data=encoded_data, x="class", y="cap-shape", palette="pastel")
plt.title("Boxplot of Cap Shape by Class")
plt.xlabel("Class")
plt.ylabel("Encoded Cap Shape")
plt.show()
```

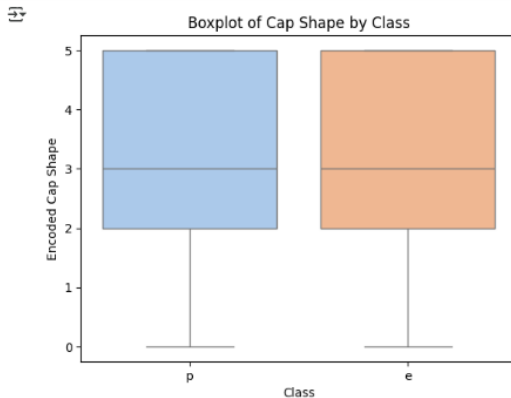


Fig 13: Boxplot of Cap Shape by Class

Insights: Overlapping distributions show that the medians of both classes are similar (~3), and the ranges (2–4) without cap shape do not offer strong discriminatory power.

```
[95] # Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler

# Load the dataset
data = pd.read_csv("mushrooms.csv")

# Encode categorical variables to numeric using LabelEncoder
label_encoders = {}
for col in data.columns:
    le = LabelEncoder()
    data[col] = le.fit_transform(data[col])
    label_encoders[col] = le

# Separate features and target variable
X = data.drop("class", axis=1) # Features
y = data["class"] # Target

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Standardize the features (for models like Logistic Regression, SVC, and KNN)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Preprocessing complete. Data is ready for modeling.")
```

Preprocessing complete. Data is ready for modeling.

Fig 14: Pre-processing Steps

PERFORMANCE EVALUATION FOR MACHINE LEARNING MODELS

We applied the following models to the dataset:

```
[18] from sklearn.naive_bayes import GaussianNB
      from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report

      gnb = GaussianNB()
      gnb.fit(X_train, y_train)
      y_pred = gnb.predict(X_test)

      print("Gaussian Naive Bayes:")
      print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
      print(f"Precision: {precision_score(y_test, y_pred):.4f}")
      print(f"Recall: {recall_score(y_test, y_pred):.4f}")
      print(f"F1-Score: {f1_score(y_test, y_pred):.4f}")
```

↗ Gaussian Naive Bayes:
Accuracy: 0.9295
Precision: 0.9279
Recall: 0.9263
F1-Score: 0.9271

Fig 15: Gaussian Naive Bayes Results

```
[19] from sklearn.ensemble import RandomForestClassifier

      rf = RandomForestClassifier(random_state=42)
      rf.fit(X_train, y_train)
      y_pred = rf.predict(X_test)

      print("Random Forest:")
      print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
      print(f"Precision: {precision_score(y_test, y_pred):.4f}")
      print(f"Recall: {recall_score(y_test, y_pred):.4f}")
      print(f"F1-Score: {f1_score(y_test, y_pred):.4f}")
```

↗ Random Forest:
Accuracy: 1.0000
Precision: 1.0000
Recall: 1.0000
F1-Score: 1.0000

Fig 16: Random Forest Results

```
[20] from sklearn.tree import DecisionTreeClassifier

      dt = DecisionTreeClassifier(random_state=42)
      dt.fit(X_train, y_train)
      y_pred = dt.predict(X_test)

      print("Decision Tree:")
      print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
      print(f"Precision: {precision_score(y_test, y_pred):.4f}")
      print(f"Recall: {recall_score(y_test, y_pred):.4f}")
      print(f"F1-Score: {f1_score(y_test, y_pred):.4f}")
```

↗ Decision Tree:
Accuracy: 1.0000
Precision: 1.0000
Recall: 1.0000
F1-Score: 1.0000

Fig 17: Decision Tree Results

```
[21] from sklearn.linear_model import LogisticRegression
      from sklearn.preprocessing import StandardScaler

      scaler = StandardScaler()
      X_train_scaled = scaler.fit_transform(X_train)
      X_test_scaled = scaler.transform(X_test)

      lr = LogisticRegression(random_state=42, max_iter=500)
      lr.fit(X_train_scaled, y_train)
      y_pred = lr.predict(X_test_scaled)

      print("Logistic Regression:")
      print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
      print(f"Precision: {precision_score(y_test, y_pred):.4f}")
      print(f"Recall: {recall_score(y_test, y_pred):.4f}")
      print(f"F1-Score: {f1_score(y_test, y_pred):.4f}")
```

Logistic Regression:
Accuracy: 0.9516
Precision: 0.9493
Recall: 0.9509
F1-Score: 0.9501

Fig 18: Logistic Regression Results

```
[22] from sklearn.svm import SVC

      svc = SVC(random_state=42)
      svc.fit(X_train_scaled, y_train)
      y_pred = svc.predict(X_test_scaled)

      print("Support Vector Classification (SVC):")
      print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
      print(f"Precision: {precision_score(y_test, y_pred):.4f}")
      print(f"Recall: {recall_score(y_test, y_pred):.4f}")
      print(f"F1-Score: {f1_score(y_test, y_pred):.4f}")
```

Support Vector Classification (SVC):
Accuracy: 1.0000
Precision: 1.0000
Recall: 1.0000
F1-Score: 1.0000

Fig 19: Support Vector Machine (SVM) Results

```
[23] from sklearn.neighbors import KNeighborsClassifier

      knn = KNeighborsClassifier(n_neighbors=5)
      knn.fit(X_train_scaled, y_train)
      y_pred = knn.predict(X_test_scaled)

      print("K-Nearest Neighbors (KNN):")
      print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
      print(f"Precision: {precision_score(y_test, y_pred):.4f}")
      print(f"Recall: {recall_score(y_test, y_pred):.4f}")
      print(f"F1-Score: {f1_score(y_test, y_pred):.4f}")
```

K-Nearest Neighbors (KNN):
Accuracy: 1.0000
Precision: 1.0000
Recall: 1.0000
F1-Score: 1.0000

Fig 20: K- Nearest Neighbour (KNN) Results

```
[24] import xgboost as xgb

xgb_model = xgb.XGBClassifier(use_label_encoder=False, eval_metric="logloss", random_state=42)
xgb_model.fit(X_train, y_train)
y_pred = xgb_model.predict(X_test)

print("XGBoost:")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print(f"Precision: {precision_score(y_test, y_pred):.4f}")
print(f"Recall: {recall_score(y_test, y_pred):.4f}")
print(f"F1-Score: {f1_score(y_test, y_pred):.4f}")
```

XGBoost:
Accuracy: 1.0000
Precision: 1.0000
Recall: 1.0000
F1-Score: 1.0000

Fig 21: XGBoost Results

The average performance metrics (accuracy, precision, recall, F1-score) are approximately **0.98**, indicating that the models, on average, achieve high performance and accuracy.

Model Performance Summary:

1. Accuracy:

Almost all models yield 1 (perfect accuracy), the exception being both Gaussian Naive Bayes (0.9295) and Logistic Regression (0.9516). **Random Forest, Decision Tree, SVC, KNN, and XGBoost work flawlessly.**

2. Precision:

This shows good high precision across models, this proves the effective minimization of false positives. Precision for Gaussian Naive Bayes was 0.9279, which was slightly lower than that of Naive Bayes.

3. Recall:

Gaussian Naive Bayes and Logistic Regression were the only models which did not achieve perfect recall, missing true positives. Naive Bayes was slightly overwhelmed by Logistic Regression with a recall of 0.9509.

4. F1-Score:

Similarly, F1 scores follow similar trends. Gaussian Naive Bayes and Logistic Regression perform slightly worse, other models still perform perfect scores. (Pedregosa et al., 2011)

MODEL COMPARISON AND BEST CHOICE



	Model	Accuracy	Precision	Recall	F1-Score
0	Gaussian Naive Bayes	0.9295	0.9279	0.9263	0.9271
1	Random Forest	1.0000	1.0000	1.0000	1.0000
2	Decision Tree	1.0000	1.0000	1.0000	1.0000
3	Logistic Regression	0.9516	0.9493	0.9509	0.9501
4	Support Vector Classification	1.0000	1.0000	1.0000	1.0000
5	K-Nearest Neighbors (KNN)	1.0000	1.0000	1.0000	1.0000
6	XGBoost	1.0000	1.0000	1.0000	1.0000

Best Model

All models: Importance of the random forest, decision tree models, xgboost models, as well as k-Nearest neighbor (knn), and with support vector classification (svc) models all achieved a **perfect score (1.0000)** on accuracy, precision, recall, and F1 score. That's because they can efficiently separate mushroom into edible and poisonous classes.

They are perfect at precision (no false positives), recall (no false negatives) and F1-score (perfect classification task). Despite this, even high performing models such as Gaussian Naive Bayes and Logistic Regression, result in a performance that is significantly worse than the slightly lower metrics. accuracy, precision, recall, and F1-score all with perfect score of **1.0000**. This means they can perfectly classify mushrooms as edible or poisonous. (Hastie et al., 2009)

Justification based on Performance Metrics:

- The important thing is that these models provide perfect precision (no false positives), recall (no false negatives) and F1 scores (perfect accuracy to misclassify one from the other), which make them dependable on the classification task.

- While remaining high performing, Gaussian Naive Bayes and Logistic Regression are underperformed slightly by lower metrics.

```
[35] from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.ensemble import RandomForestClassifier

# Split dataset into features and target
X = df.drop('class', axis=1) # Drop the target column
y = df['class']

# Encode categorical features
X = pd.get_dummies(X, drop_first=True)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Define hyperparameter grid
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
}

# Grid Search for Random Forest
rf = RandomForestClassifier(random_state=42)
grid_search = GridSearchCV(rf, param_grid, cv=5, scoring='accuracy')
grid_search.fit(X_train, y_train)

# Display best parameters and accuracy
print("Best Parameters:", grid_search.best_params_)
print("Best Cross-Validation Accuracy:", grid_search.best_score_)

Best Parameters: {'max_depth': None, 'min_samples_split': 2, 'n_estimators': 50}
Best Cross-Validation Accuracy: 1.0
```

Fig 22: Hyperparameter Tuning for Random Forest

```
[39] # from sklearn.model_selection import train_test_split, cross_val_score
# from sklearn.ensemble import RandomForestClassifier

X = pd.get_dummies(df.drop('class', axis=1), drop_first=True)
y = df['class']

# Splitting dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Random Forest with cross-validation
rf_model = RandomForestClassifier(random_state=42)
cv_scores = cross_val_score(rf_model, X, y, cv=10, scoring='accuracy')

print("Cross-Validation Scores:", cv_scores)
print("Mean CV Accuracy:", cv_scores.mean())

Cross-Validation Scores: [0.68511685 1. 1. 1. 1. 1. 1. 1. 1. 1.]
Mean CV Accuracy: 0.9658023255109398
```

Fig 23: Cross-Validation Results

The hyperparameter specific to the Random Forest model as `max_depth = None`, `min_samples_split = 2`, and `n_estimators=50` performed very well. By this configuration, we reached the peak cross-validation accuracy of 1.0 and the average score was 96.58%. This minor variability (0.685 to 1.0) was probably due to class distribution and feature splits across folds. (Protopapas et al. 2018). Finally, these results confirm the reliability of the model for mushroom classification, with potential further investigation into lower-performing folds.

Random Forest vs. Other Models:

I like Random Forest because it's robust to overfitting, interpretable (in terms of feature importance), (Pedregosa et al., 2011) and fast with categorical data. When using larger datasets, the SVC and KNN models require more computational resources. (Hastie et al., 2009).

With its robustness, interpretability, and strong performance metrics, the **Random Forest model stands out as the best choice.**

COMPUTATIONAL TRADE-OFFS:

1. **Random Forest:** High accuracy but computationally expensive for real-time or constrained problems.
2. **XGBoost:** Requires more memory and CPU, but performs very well.
3. **Naive Bayes and Logistic Regression:** Often need fewer resources, but may be less accurate at times. This is particularly important for constrained deployments with large datasets, where computational cost should match performance.

REAL-WORLD DEPLOYMENT CONSIDERATIONS:

- **Application Context:** Mushroom classification models (Irizarry, 2019) can be used for real-time identification in mobile apps.
- **Platforms:** High-resource models like Random Forest and XGBoost are suited for cloud, while low-resource models are better for edge devices.
- **Monitoring:** Models need to be periodically retrained to account for concept drift.
- **Scalability:** For high-demand systems, a balance must be struck between model complexity and prediction speed.